

Missing Person Face Search System

This document explains **everything completed so far**, step by step, so you (or anyone else) can set up the system from scratch without confusion.

The system is designed using **industry-grade architecture**:

- **Qdrant** → vector database (face embeddings)
- **PostgreSQL** → metadata database (name, state, police station, etc.)
- **InsightFace** → face detection & embedding
- **Docker** → run databases cleanly

Project Structure (Current)

```
final_budddy/
├── final_images/          # FinalPersonId.jpg images (1 face per person)
├── faces_found.csv       # Metadata of persons with valid faces
├── faces_not_found.csv   # Persons with no usable face
├── create_qdrant_collection.py
├── ingest_embeddings.py
├── progress_checkpoint.csv # Resume-safe checkpoint
├── qdrant_storage/        # Qdrant persistent data (Docker volume)
├── pgdata/                # PostgreSQL persistent data (Docker volume)
└── README.md              # This file
```

System Architecture (High Level)

```
Image (.jpg)
  ↓
InsightFace (buffalo_s)
  ↓
Face Embedding (512-D)
  ↓
Qdrant Vector DB — FinalPersonId — PostgreSQL Metadata DB
```

- **Qdrant** stores only vectors + `FinalPersonId`
- **PostgreSQL** stores all human-readable metadata

Prerequisites

Python Environment

```
python --version # Python 3.10+ recommended
```

Create virtual environment:

```
python -m venv .venv  
source .venv/bin/activate
```

Install dependencies:

```
pip install insightface qdrant-client psycopg2-binary pillow numpy tqdm  
requests
```

Docker Setup

Verify Docker Installation

```
docker --version  
docker ps
```

If Docker is not running (macOS):

```
open -a Docker
```

Wait until Docker Desktop shows “**Docker is running**”.

Start Qdrant (Vector Database)

Run Qdrant Container

```
docker run -d  
--name qdrant  
-p 6333:6333  
-p 6334:6334  
-v $(pwd)/qdrant_storage:/qdrant/storage  
qdrant/qdrant
```

Verify Qdrant

```
docker ps
```

Open dashboard in browser:

```
http://localhost:6333/dashboard
```

Start PostgreSQL (Metadata Database)

Run PostgreSQL Container

```
docker run -d
--name postgres
-e POSTGRES_DB=faces_db
-e POSTGRES_USER=postgres
-e POSTGRES_PASSWORD=postgres
-p 5432:5432
-v $(pwd)/pgdata:/var/lib/postgresql/data
postgres:15
```

Verify PostgreSQL

```
docker ps
```

Connect to PostgreSQL (Recommended Way)

You **do not need** `psql` **installed locally**.

```
docker exec -it postgres psql -U postgres -d faces_db
```

Password:

```
postgres
```

Inside psql:

```
\dt
```

Exit:

```
\q
```

👉 Create Qdrant Collection

💡 Run Collection Creation Script

```
python create_qdrant_collection.py
```

Expected output:

✓ Qdrant collection created with payload index

(or)

ℹ️ Qdrant collection already exists

Collection details: - Name: `missing_person_faces` - Vector size: `512` - Distance: `COSINE`

🧠 Dataset Preparation (Already Done)

What exists already

- Face-filtered images in `final_images/`
- One image per `FinalPersonId`
- `faces_found.csv` contains:

`FinalPersonId, Name, Sex, BirthYear, State, District,
PoliceStation, TracingStatus, ImageFile`

😢 Embedding + Metadata Ingestion

💡 Run Ingestion Script

```
python ingest_embeddings.py
```

What this script does:

- Reads `faces_found.csv`
- Loads `final_images/FP_xxx.jpg`
- Generates face embedding (`buffalo_s`)
- Inserts embedding into **Qdrant**
- Inserts metadata into **PostgreSQL**
- Skips already-inserted `FinalPersonId`

This script is **resume-safe**.

👉 Verify Data After Ingestion

Qdrant

```
curl http://localhost:6333/collections
```

PostgreSQL

```
docker exec -it postgres psql -U postgres -d faces_db
```

```
SELECT COUNT(*) FROM persons;  
SELECT * FROM persons LIMIT 5;
```

👉 Current Status Checklist

Component	Status
Image dataset	✓ Ready
Face filtering	✓ Done
Qdrant running	✓
PostgreSQL running	✓
Vector ingestion	✓
Metadata ingestion	✓

😢 Next Steps (Future)

- ⚡ FastAPI face search API
- 📊 Admin dashboards

- Threshold tuning
 -  Authentication
 -  GPU acceleration
-

Summary

You now have a **real, production-style face search backend**:

- Clean separation of vector search and metadata
- Resume-safe pipelines
- Scalable architecture
- Ready for APIs and UI

This is **not a demo system** — it's how serious AI backends are built.